

Jev vs. Other AI Models

Built for **10k-analyst** — an AI analyst for SEC 10-K filings
Jev (TypeSafe) used for bounded judgments · Claude used for cited answers

Jev (TypeSafe System One)

Typed judgments: pick one, score one, answer yes/no

General-purpose LLM (e.g. ChatGPT)

Open-ended text generation and reasoning

Traditional ML Classifier

A model trained on labeled examples for one fixed task


Speed, cost & setup

Lower setup effort and higher speed/cost-efficiency score better

Speed per call



Cost efficiency



Zero-shot readiness (less setup needed)




Illustrative relative rating, 1–5 scale — not a benchmarked study.


Reliability & capability

Based on each approach's documented characteristics

Confidence calibration




Numeric reliability



Flexibility to new tasks



Explainability



Illustrative relative rating, 1–5 scale — not a benchmarked study.

Pros & Cons by Approach

Jev

TypeSafe System One

PROS

- Cheap and fast enough to run on every paragraph
- Calibrated confidence on every answer
- Structurally cannot answer outside the options given
- No training data needed to start

CONS

- Only bounded judgments — no open-ended writing
- Weak at numbers and literal detail by design
- Newer ecosystem, fewer public benchmarks

BEST FIT

High-volume bounded decisions embedded directly in app logic — routing, relevance filtering, and verification gates, like the ones in this project's retrieval and groundedness checks.

General-purpose LLM

e.g. ChatGPT, Claude, Gemini

PROS

- Handles nearly any task via prompting
- Produces fluent, cited, human-readable prose
- Can reason step by step and explain itself

CONS

- Slower and costlier per call, especially at volume
- Confidence must be prompted for and is often unreliable
- Can hallucinate facts, figures, or sources

BEST FIT

Open-ended writing, synthesis, and reasoning a human analyst would otherwise do — drafting the cited answer itself, as Claude does in this project.

Traditional ML Classifier

e.g. a fine-tuned BERT or logistic regression model

PROS

- Extremely fast and cheap once deployed
- Confidence calibration is a mature, solved practice
- Predictable, auditable behavior at scale

CONS

- Needs labeled training data and a training pipeline
- Locked to the task it was trained for
- Expensive to adapt when requirements change

BEST FIT


Stable, well-defined, high-scale tasks with plentiful labeled data and infrequently changing requirements.

Why Jev?

Jev is the clear choice for the bounded judgments this project relies on: tagging risk topics, filtering relevant passages, and gating groundedness. It returns a **calibrated confidence score** on every single answer, something general-purpose LLMs don't do natively. It runs **fast** and **cheap** enough to evaluate every paragraph in a 10-K filing without a second thought, and because its output is always one of a **fixed, typed set of answers**, it can **never drift outside the options it's given**, a structural guarantee free-text generation simply can't match. Across the dimensions that matter most for this class of problem, **Jev comes out on top**.

Where Jev leads

Same ratings shown on page 1, reproduced here as a vertical comparison



Illustrative relative rating, 1–5 scale. Not a benchmarked study. Jev reaches the top of its scale on every dimension shown.